

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee64](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee64)

Enter your authorization code:  
.....  
Mounted at /content/drive

```
1 !nvidia-smi
```

Mon May 11 26:36 2020

+-----+-----+-----+-----+-----+									
NVIDIA-SMI		440.82		Driver Version: 418.67			CUDA Version: 10.1		
+-----+-----+-----+-----+-----+									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr. ECC		
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.		
=====									
0	Tesla K80	Off		00000000:00:04.0 Off		0			
N/A	32C	P8	29W / 149W	0MiB / 11441MiB		0%	Default		
+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+									
Processes:								GPU Memory	
GPU	PID	Type	Process name				Usage		
=====									
No running processes found									
+-----+-----+-----+-----+-----+									

Dataset prepared for Pytorch Pipeline

```
1 !gdown --id 1cV2bDv5h0_wS7fANNdQPWM40hE3PyWls
2 !unzip springfield_dataset.zip
```



```
inflating: springfield_dataset/val/sideshow_bob/pic_0173.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0174.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0179.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0194.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0234.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0243.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0249.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0264.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0269.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0271.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0275.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0283.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0287.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0304.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0312.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0331.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0371.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0378.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0400.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0413.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0424.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0430.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0442.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0466.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0478.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0486.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0498.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0502.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0513.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0527.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0537.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0538.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0543.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0549.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0552.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0557.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0565.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0573.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0578.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0580.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0600.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0617.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0633.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0635.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0646.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0658.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0661.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0663.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0672.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0676.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0682.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0686.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0687.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0716.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0720.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0731.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0742.jpg
inflating: springfield_dataset/val/sideshow_bob/pic_0747.jpg
```

```
1 from __future__ import print_function, division
2
3 import torch
4 import torch.nn as nn
5 import torch.optim as optim
6 from torch.optim import lr_scheduler
7 import numpy as np
8 import torchvision
9 from torchvision import datasets, models, transforms
10 import matplotlib.pyplot as plt
11 import time
12 import os
13 import copy
14
15 plt.ion()    # interactive mode
```

```
inflating: springfield_dataset/val/sideshow_mel/pic_0006.jpg
```

```
1 !ls springfield_dataset/
```

```
➤ characters_illustration.png  sample_submission.csv  train
  make_val.py                  testset              val
  inflating: springfield_dataset/val/snake_tailbird/pic_0006.png
```

Create Training and Validation sets.

Prepare input images.

```
inflating: springfield_dataset/val/troy_mcclure/pic_0006.png
```

```
1 data_transforms = {
2     'train': transforms.Compose([
```

```
1 print(dataloaders)
2 print(dataset_sizes)
3 print(class_names)
```

```

1 def imshow(inp, title=None):
2     """Imshow for Tensor."""
3     inp = inp.numpy().transpose((1, 2, 0))
4     mean = np.array([0.485, 0.456, 0.406])
5     std = np.array([0.229, 0.224, 0.225])
6     inp = std * inp + mean
7     inp = np.clip(inp, 0, 1)
8     plt.imshow(inp)
9     if title is not None:
10         plt.title(title)
11     plt.pause(0.001) # pause a bit so that plots are updated
12
13
14 # Get a batch of training data
15 inputs, classes = next(iter(dataloaders['train']))
16
17 # Make a grid from batch
18 out = torchvision.utils.make_grid(inputs)
19
20 imshow(out, title=[class_names[x] for x in classes])

```


`['charles_montgomery_burns', 'homer_simpson', 'homer_simpson', 'principal_skinner', 'comic_book_guy', 'moe_szyzlak', 'homer_simpson', 'ned_flanders', 'krusty_the_clown', 'maggie_simpson', 'ideshow_bob', 'moe_szyzlak', 'ned_flanders', 'kent_brockman', 'gil', 'homer_simpson']`



```

1 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
2     since = time.time()
3
4     best_model_wts = copy.deepcopy(model.state_dict())
5     best_acc = 0.0
6
7     for epoch in range(num_epochs):
8         print('Epoch {}/{}'.format(epoch, num_epochs - 1))
9         print('-' * 10)
10
11        # Each epoch has a training and validation phase
12        for phase in ['train', 'val']:
13            if phase == 'train':
14                model.train()  # Set model to training mode
15            else:
16                model.eval()    # Set model to evaluate mode

```

```

17
18     running_loss = 0.0
19     running_corrects = 0
20
21     # Iterate over data.
22     for inputs, labels in dataloaders[phase]:
23         inputs = inputs.to(device)
24         labels = labels.to(device)
25
26         # zero the parameter gradients
27         optimizer.zero_grad()
28
29         # forward
30         # track history if only in train
31         with torch.set_grad_enabled(phase == 'train'):
32             outputs = model(inputs)
33             _, preds = torch.max(outputs, 1)
34             loss = criterion(outputs, labels)
35
36         # backward + optimize only if in training phase
37         if phase == 'train':
38             loss.backward()
39             optimizer.step()
40
41         # statistics
42         running_loss += loss.item() * inputs.size(0)
43         running_corrects += torch.sum(preds == labels.data)
44     if phase == 'train':
45         scheduler.step()
46
47     epoch_loss = running_loss / dataset_sizes[phase]
48     epoch_acc = running_corrects.double() / dataset_sizes[phase]
49
50     print('{} Loss: {:.4f} Acc: {:.4f}'.format(
51         phase, epoch_loss, epoch_acc))
52
53     # deep copy the model
54     if phase == 'val' and epoch_acc > best_acc:
55         best_acc = epoch_acc
56         best_model_wts = copy.deepcopy(model.state_dict())
57
58     print()
59
60     time_elapsed = time.time() - since
61     print('Training complete in {:.0f}m {:.0f}s'.format(
62         time_elapsed // 60, time_elapsed % 60))
63     print('Best val Acc: {:.4f}'.format(best_acc))
64
65     # load best model weights
66     model.load_state_dict(best_model_wts)
67     return model

```

```

1 def visualize_model(model, num_images=6):
2     was_training = model.training
3     model.eval()
4     images_so_far = 0
5     fig = plt.figure()
6
7     with torch.no_grad():
8         for i, (inputs, labels) in enumerate(dataloaders['val']):
9             inputs = inputs.to(device)
10            labels = labels.to(device)
11
12            outputs = model(inputs)
13            _, preds = torch.max(outputs, 1)
14
15            for j in range(inputs.size()[0]):
16                images_so_far += 1
17                ax = plt.subplot(num_images//2, 2, images_so_far)
18                ax.axis('off')
19                ax.set_title('predicted: {}'.format(class_names[preds[j]]))
20                imshow(inputs.cpu().data[j])
21
22            if images_so_far == num_images:
23                model.train(mode=was_training)
24                return
25     model.train(mode=was_training)

```

TRAIN

Optional step, you can **skip** because the next step contains trained weights.

Transfer learning was used.

Pretrained **resnext50\_32x4d** has been trained for 10 epochs. The best parameters were saved in the model.

Parameters will be restored in the next step.

```
1 # model_ft = models.resnet18(pretrained=True)
2 import torchvision.models as models
3 model_ft = models.resnext50_32x4d(pretrained=True)
4 num_fts = model_ft.fc.in_features
5 # Here the size of each output sample is set to 2.
6 # Alternatively, it can be generalized to nn.Linear(num_fts, len(class_names)).
7 model_ft.fc = nn.Linear(num_fts, 42)
8
9 model_ft = model_ft.to(device)
10
11 criterion = nn.CrossEntropyLoss()
12
13 # Observe that all parameters are being optimized
14 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
15
16 # Decay LR by a factor of 0.1 every 7 epochs
17 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)
```

📁 Downloading: "[https://download.pytorch.org/models/resnext50\\_32x4d-7cdf4587.pth](https://download.pytorch.org/models/resnext50_32x4d-7cdf4587.pth)" to /root/.cache/torch/checkpoints/resnext50\_32x4d-7cdf4587.pth 100% 95.8M/95.8M [03:09<00:00, 530kB/s]

```
1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
2                         num_epochs=10)
```

**Recover the weight** and initialize the new resnext50\_32x4d model with them.

```
1 # Download parameters
2 !gdown --id 1RB9RbRt1_VlLTuKLaLHtTEXhyn_NRoo7
```

📁 Downloading...  
From: [https://drive.google.com/uc?id=1RB9RbRt1\\_VlLTuKLaLHtTEXhyn\\_NRoo7](https://drive.google.com/uc?id=1RB9RbRt1_VlLTuKLaLHtTEXhyn_NRoo7)  
To: /content/resnext.pt  
92.6MB [00:00, 223MB/s]

```
1 # Initialize a new model with the best parameters restored
2 model_ft = models.resnext50_32x4d()
3 num_fts = model_ft.fc.in_features
4 model_ft.fc = nn.Linear(num_fts, 42)
5 model_ft = model_ft.to(device)
6
7 path = F"/content/resnext.pt"
8 model_ft.load_state_dict(torch.load(path))
9 model_ft.eval()
```

📁

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  )
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer3): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(512, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=32, bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
)

```

```

        (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (3): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (4): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (5): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=32, bias=False)
    (bn2): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (downsample): Sequential(
    (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
  (1): Bottleneck(

```

```

1 criterion = nn.CrossEntropyLoss()
2
3 # Observe that all parameters are being optimized
4 optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
5
6 # Decay LR by a factor of 0.1 every 7 epochs
7 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

```

Try to improve current results.

This step **didn't** bring improvement. Skip

```

    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
2                           num_epochs=25)
3

```



```
Epoch 0/24
-----
train Loss: 0.3186 Acc: 0.9148
val Loss: 0.0428 Acc: 0.9899
```

```
Epoch 1/24
-----
train Loss: 0.3194 Acc: 0.9140
val Loss: 0.0545 Acc: 0.9851
```

```
Epoch 2/24
-----
train Loss: 0.2957 Acc: 0.9226
val Loss: 0.0402 Acc: 0.9913
```

```
Epoch 3/24
-----
train Loss: 0.3056 Acc: 0.9185
val Loss: 0.0459 Acc: 0.9899
```

```
Epoch 4/24
-----
train Loss: 0.2778 Acc: 0.9247
val Loss: 0.0543 Acc: 0.9851
```

```
Epoch 5/24
-----
train Loss: 0.2747 Acc: 0.9260
val Loss: 0.0526 Acc: 0.9865
```

```
Epoch 6/24
-----
train Loss: 0.2592 Acc: 0.9287
val Loss: 0.0462 Acc: 0.9880
```

```
Epoch 7/24
-----
train Loss: 0.2351 Acc: 0.9359
val Loss: 0.0391 Acc: 0.9894
```

```
Epoch 8/24
-----
train Loss: 0.2253 Acc: 0.9410
val Loss: 0.0391 Acc: 0.9894
```

```
Epoch 9/24
-----
```

```
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-214-cc88ea5f8bd3> in <module>()
      1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,
----> 2                                num_epochs=25)

----- 2 frames -----
/usr/local/lib/python3.6/dist-packages/torch/autograd/__init__.py in backward(tensors, grad_tensors, retain_graph, create_graph,
      98     Variable._execution_engine.run_backward(
      99         tensors, grad_tensors, retain_graph, create_graph,
--> 100         allow_unreachable=True) # allow_unreachable flag
      101
      102
```

KeyboardInterrupt:

KeyboardInterrupt:

Save the best weights. **Not necessary**

```
1 model_save_name = 'resnext.pt'
2 path = F"/content/drive/My Drive/{model_save_name}"
3 torch.save(model_ft.state_dict(), path)
```

Validation

```
1 # Iterate over data.
2 val_preds = np.array([], dtype='int32')
3 val_labels = np.array([], dtype='int32')
4 running_corrects = 0
5
6 for inputs, labels in dataloaders['val']:
7     inputs = inputs.to(device)
8     labels = labels.to(device)
9
10    outputs = model_ft(inputs)
11    _, preds = torch.max(outputs, 1)
```



```

12
13     # statistics
14     running_corrects += torch.sum(preds == labels.data)
15
16     preds = preds.cpu().numpy()
17     labels = labels.cpu().numpy()
18
19     val_preds = np.concatenate((val_preds, preds), axis=None)
20     val_labels = np.concatenate((val_labels, labels), axis=None)

```

```

1 print(val_preds)
2 print(val_labels)
3 (unique, counts) = np.unique(val_preds, return_counts=True)
4 frequencies = np.asarray((unique, counts)).T
5 print(frequencies)

```

```

[16 15  2 ...  7 37  6]
[16 15  2 ...  7 37  6]
[[ 0  91]
 [ 1   4]
 [ 2  61]
 [ 3  10]
 [ 4 134]
 [ 5   9]
 [ 6 119]
 [ 7  97]
 [ 8   4]
 [ 9  46]
[11  44]
[12   2]
[13   2]
[14  11]
[15 226]
[16  49]
[17 120]
[18  31]
[20 134]
[21  12]
[22 128]
[23   9]
[24  25]
[25 109]
[26   1]
[27 149]
[28 146]
[29  34]
[30   3]
[31   7]
[32 118]
[33   6]
[34   3]
[35   8]
[36  10]
[37  87]
[38   4]
[39   5]
[41  18]]

```

```

1 from sklearn.metrics import f1_score
2 print(f1_score(val_labels, val_preds, average='micro'))
3 print(f1_score(val_labels, val_preds, average='macro'))
4 print(f1_score(val_labels, val_preds, average='weighted'))

```

```

[0.9894026974951831
 0.9086346522934735
 0.988714637298817]

```

## Test

In order not to take risks with the final results, the pipeline DLS is used

```

1 # Download label_encoder
2 !gdown --id 15UI2ogJbAP2u04AKKj8byqqv1SPCbbGi

```

```

[Downloading...]
From: https://drive.google.com/uc?id=15UI2ogJbAP2u04AKKj8byqqv1SPCbbGi
To: /content/label_encoder.pkl
100% 4.30k/4.30k [00:00<00:00, 7.42MB/s]

```

```

1 import pickle
2 from pathlib import Path
3

```

```

4 TEST_DIR = Path('/content/springfield_dataset/testset/testset')
5 test_files = sorted(list(TEST_DIR.rglob('*.jpg')))
6 DATA_MODES = ['train', 'val', 'test']
7 RESCALE_SIZE = 224
8
9 label_encoder = pickle.load(open("label_encoder.pkl", 'rb'))

```

```
1 print(len(test_files))
```

991

```

1 from torch.utils.data import Dataset, DataLoader
2 from sklearn.preprocessing import LabelEncoder
3 from PIL import Image
4
5 class SimpsonsDataset(Dataset):
6     """
7     Датасет с картинками, который паралельно подгружает их из папок
8     производит скалирование и превращение в торчевые тензоры
9     """
10    def __init__(self, files, mode):
11        super().__init__()
12        # список файлов для загрузки
13        self.files = sorted(files)
14        # режим работы
15        self.mode = mode
16
17        if self.mode not in DATA_MODES:
18            print(f"{self.mode} is not correct; correct modes: {DATA_MODES}")
19            raise NameError
20
21        self.len_ = len(self.files)
22
23        self.label_encoder = LabelEncoder()
24
25        if self.mode != 'test':
26            self.labels = [path.parent.name for path in self.files]
27            self.label_encoder.fit(self.labels)
28
29            with open('label_encoder.pkl', 'wb') as le_dump_file:
30                pickle.dump(self.label_encoder, le_dump_file)
31
32    def __len__(self):
33        return self.len_
34
35    def load_sample(self, file):
36        image = Image.open(file)
37        image.load()
38        return image
39
40    def __getitem__(self, index):
41        # для преобразования изображений в тензоры PyTorch и нормализации входа
42        transform = transforms.Compose([
43            transforms.ToTensor(),
44            transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
45        ])
46        x = self.load_sample(self.files[index])
47        x = self._prepare_sample(x)
48        x = np.array(x / 255, dtype='float32')
49        x = transform(x)
50        if self.mode == 'test':
51            return x
52        else:
53            label = self.labels[index]
54            label_id = self.label_encoder.transform([label])
55            y = label_id.item()
56            return x, y
57
58    def _prepare_sample(self, image):
59        image = image.resize((RESCALE_SIZE, RESCALE_SIZE))
60        return np.array(image)

```

```

1 def predict(model, test_loader):
2     with torch.no_grad():
3         logits = []
4
5         for inputs in test_loader:
6             inputs = inputs.to(device)
7             model.eval()

```

```
7         model.eval()
8         outputs = model(inputs).cpu()
9         logits.append(outputs)
10
11     probs = nn.functional.softmax(torch.cat(logits), dim=-1).numpy()
12     return probs
```

Make predictions

**The final Test Result:** 0.99468

**Kaggle Nickname:** Mezga\_Alexander\_35630269

```
1 test_dataset = SimpsonsDataset(test_files, mode="test")
2 test_loader = DataLoader(test_dataset, shuffle=False, batch_size=64)
```

```
1 probs = predict(model_ft, test_loader)
```

```
1 preds = label_encoder.inverse_transform(np.argmax(probs, axis=1))
2 test_filenames = [path.name for path in test_dataset.files]
```

Create submission and save to the Drive

```
1 import pandas as pd
2 # my_submit = pd.read_csv("gdrive/My Drive/simpsons/data/labels.csv")
3 my_submit = pd.DataFrame({'Id': test_filenames, 'Expected': preds})
4 my_submit.head()
5 my_submit.to_csv('/content/drive/My Drive/DLS_Hometask/resnet_baseline_2.csv', index=False)
```

```
1
```